

# Где дешевле жить? Предсказание цен в Airbnb

Опробованные техники.

- Была сделана стандартная предобработка признаков
- построены различные визуализации
- опробован one через dummies
- применены различные типы шкалирования (MinMax, StandardScaler, Robust)
- для переменной neighborhood применено частотное кодирование, чтобы не перегружать число признаков
- применена помимо линейных моделей - random forest

Итоги:

- фича превращения long и lat в расстояние от центра манхетена до недвижимости на удивление не помогла
- заметно улучшило модель шкалирование данных, лучше всего сработало robust шкалирование
- в целом линейные модели показали примерно одинаковый результат
- random forest оказался немного лучше линейных моделей

Все итоговые метрики расположены в конце ноутбука в dataframe measured\_metrics

## Часть 1. EDA

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
pd.set_option('display.max_columns', None)

from scipy import stats
from scipy.stats import norm
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import RobustScaler
from sklearn.preprocessing import MinMaxScaler
from sklearn import metrics
from sklearn.linear_model import LinearRegression, RidgeCV, LassoCV, ElasticNetCV, Lasso

%matplotlib inline
plt.rcParams["figure.figsize"] = (12,8)

np.random.seed(49)
```

для отображения графиков без скроллов


```
In [2]: %%javascript
IPython.OutputArea.auto_scroll_threshold = 9999;
```

```
In [3]: data = pd.read_csv('AB_NYC_2019.csv')
```

In [4]: `data.head()`

Out[4]:

	id	name	host_id	host_name	neighbourhood_group	neighbourhood	latitude
0	2539	Clean & quiet apt home by the park	2787	John	Brooklyn	Kensington	40.64749
1	2595	Skylit Midtown Castle	2845	Jennifer	Manhattan	Midtown	40.75362
2	3647	THE VILLAGE OF HARLEM....NEW YORK !	4632	Elisabeth	Manhattan	Harlem	40.80902
3	3831	Cozy Entire Floor of Brownstone	4869	LisaRoxanne	Brooklyn	Clinton Hill	40.68514
4	5022	Entire Apt: Spacious Studio/Loft by central park	7192	Laura	Manhattan	East Harlem	40.79851



In [5]: `data.columns`

Out[5]: Index(['id', 'name', 'host\_id', 'host\_name', 'neighbourhood\_group',  
'neighbourhood', 'latitude', 'longitude', 'room\_type', 'price',  
'minimum\_nights', 'number\_of\_reviews', 'last\_review',  
'reviews\_per\_month', 'calculated\_host\_listings\_count',  
'availability\_365'],  
dtype='object')

In [6]: `use_cols = data.columns[4:12].to_list() + data.columns[13:].to_list()`

In [7]: `use_cols`

Out[7]: ['neighbourhood\_group',  
'neighbourhood',  
'latitude',  
'longitude',  
'room\_type',  
'price',  
'minimum\_nights',  
'number\_of\_reviews',  
'reviews\_per\_month',  
'calculated\_host\_listings\_count',  
'availability\_365']

In [8]: `data = pd.read_csv('AB_NYC_2019.csv', usecols=use_cols)`

In [9]: data.head()

Out[9]:

	neighbourhood_group	neighbourhood	latitude	longitude	room_type	price	minimum_nights
0	Brooklyn	Kensington	40.64749	-73.97237	Private room	149	1
1	Manhattan	Midtown	40.75362	-73.98377	Entire home/apt	225	1
2	Manhattan	Harlem	40.80902	-73.94190	Private room	150	3
3	Brooklyn	Clinton Hill	40.68514	-73.95976	Entire home/apt	89	1
4	Manhattan	East Harlem	40.79851	-73.94399	Entire home/apt	80	10

In [10]: data.describe()

Out[10]:

	latitude	longitude	price	minimum_nights	number_of_reviews	reviews_per_month
count	48895.000000	48895.000000	48895.000000	48895.000000	48895.000000	38843.000000
mean	40.728949	-73.952170	152.720687	7.029962	23.274466	0.815172
std	0.054530	0.046157	240.154170	20.510550	44.550582	1.845057
min	40.499790	-74.244420	0.000000	1.000000	0.000000	0.000000
25%	40.690100	-73.983070	69.000000	1.000000	1.000000	0.000000
50%	40.723070	-73.955680	106.000000	3.000000	5.000000	0.000000
75%	40.763115	-73.936275	175.000000	5.000000	24.000000	0.000000
max	40.913060	-73.712990	10000.000000	1250.000000	629.000000	16.340000

In [11]: data.info()


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48895 entries, 0 to 48894
Data columns (total 11 columns):
neighbourhood_group    48895 non-null object
neighbourhood          48895 non-null object
latitude               48895 non-null float64
longitude              48895 non-null float64
room_type              48895 non-null object
price                  48895 non-null int64
minimum_nights         48895 non-null int64
number_of_reviews      48895 non-null int64
reviews_per_month      38843 non-null float64
calculated_host_listings_count  48895 non-null int64
availability_365       48895 non-null int64
dtypes: float64(3), int64(5), object(3)
memory usage: 4.1+ MB
```

```
In [12]: data.dtypes
```

```
Out[12]: neighbourhood_group    object  
neighbourhood                  object  
latitude                      float64  
longitude                     float64  
room_type                     object  
price                         int64  
minimum_nights                int64  
number_of_reviews              int64  
reviews_per_month              float64  
calculated_host_listings_count int64  
availability_365               int64  
dtype: object
```

```
In [13]: data.price.hist(bins = 100)
```

```
Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x20273c26088>
```




целевая переменная распределена ненормально, пока оставим ее как есть, чтобы в дальнейшем сравнить метрики до и после предобработки.

карта района:

```
In [14]: sns.scatterplot(data.latitude, data.longitude, hue=data.neighbourhood_group)
```


```
Out[14]: <matplotlib.axes._subplots.AxesSubplot at 0x202746101c8>
```



тип комнаты в зависимости от района

```
In [15]: sns.scatterplot(data.latitude, data.longitude, hue=data.room_type)
```

```
Out[15]: <matplotlib.axes._subplots.AxesSubplot at 0x20274f36a88>
```




```
In [16]: data.neighbourhood_group.value_counts()
```

```
Out[16]: Manhattan      21661  
Brooklyn      20104  
Queens        5666  
Bronx         1091  
Staten Island    373  
Name: neighbourhood_group, dtype: int64
```

```
In [17]: sns.countplot(x = data.neighbourhood_group)
```

```
Out[17]: <matplotlib.axes._subplots.AxesSubplot at 0x20274598688>
```



проверка на дубликаты

```
In [18]: data.duplicated().sum()
```

```
Out[18]: 0
```

проверка на пропуски

```
In [19]: data.isnull().sum()
```

```
Out[19]: neighbourhood_group      0
neighbourhood                    0
latitude                        0
longitude                       0
room_type                       0
price                           0
minimum_nights                  0
number_of_reviews               0
reviews_per_month              10052
calculated_host_listings_count  0
availability_365                0
dtype: int64
```



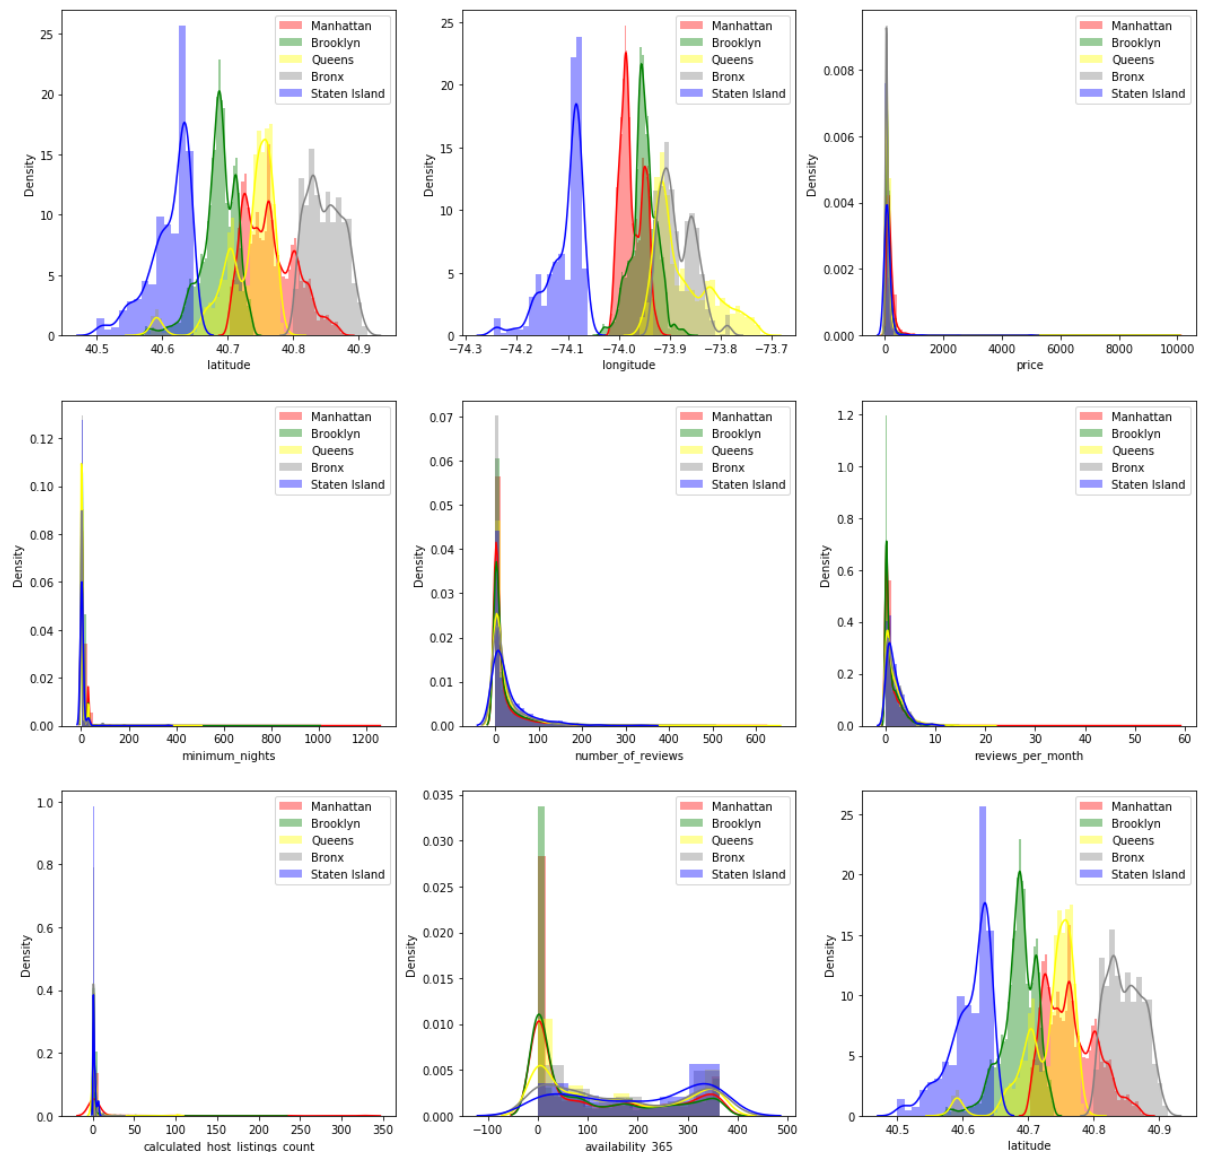
гистограммы распределения признаков

```
In [20]: int_columns = ['latitude', 'longitude', 'price', 'minimum_nights', 'number_of_reviews', 'reviews_per_month', 'calculated_host_listings_count', 'availability_365']
```


```

In [21]: nrows = 3
ncols = 3
fig, ax = plt.subplots(nrows, ncols, figsize=(ncols*6, nrows*6))
a = -8
for i in range(nrows):
    for j in range(ncols):
        ax1 = sns.distplot(data[data['neighbourhood_group'] == 'Manhattan'][int_
columns[a]], ax = ax[i,j], color = 'red', label = 'Manhattan').legend()
        ax2 = sns.distplot(data[data['neighbourhood_group'] == 'Brooklyn'][int_c
olumns[a]], ax = ax[i,j], color = 'green', label = 'Brooklyn').legend()
        ax3 = sns.distplot(data[data['neighbourhood_group'] == 'Queens'][int_col
ums[a]], ax = ax[i,j], color = 'yellow', label = 'Queens').legend()
        ax4 = sns.distplot(data[data['neighbourhood_group'] == 'Bronx'][int_colu
ms[a]], ax = ax[i,j], color = 'gray', label = 'Bronx').legend()
        ax5 = sns.distplot(data[data['neighbourhood_group'] == 'Staten Island'][
int_columns[a]], ax = ax[i,j], color = 'blue', label = 'Staten Island').legend
()
        a +=1
plt.show()

```

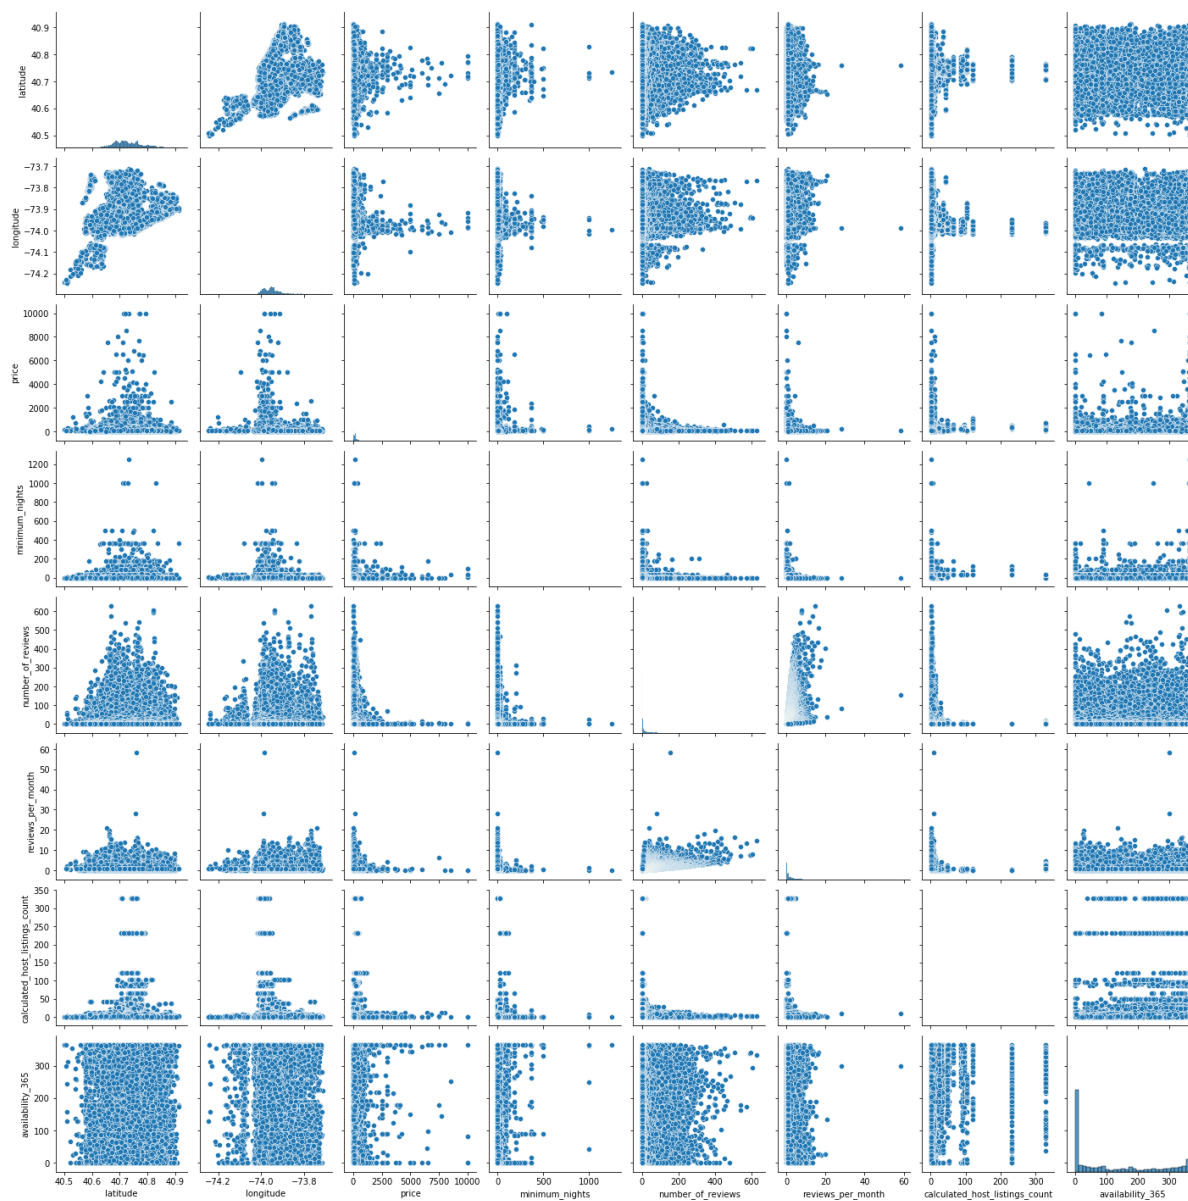


```
In [22]: corr = data.corr()  
sns.heatmap(corr, annot=True, fmt='.2f', mask = np.triu(np.ones_like(corr, dtype=bool))));
```



```
In [23]: sns.pairplot(data)
```

```
Out[23]: <seaborn.axisgrid.PairGrid at 0x2027443c608>
```



## Часть 2. Preprocessing & Feature Engineering


neighbourhood\_group

```
In [24]: data.neighbourhood_group.value_counts()
```

```
Out[24]: Manhattan      21661  
Brooklyn      20104  
Queens        5666  
Bronx         1091  
Staten Island   373  
Name: neighbourhood_group, dtype: int64
```

```
In [25]: data.neighbourhood_group.hist()
```

```
Out[25]: <matplotlib.axes._subplots.AxesSubplot at 0x2020150e208>
```



применим one-hot encoding

```
In [26]: data = pd.get_dummies(data, columns=['neighbourhood_group'], prefix='neighbourhood_group', drop_first=True)
```


**neighbourhood**

```
In [27]: data.neighbourhood.value_counts()
```

```
Out[27]: Williamsburg      3920
Bedford-Stuyvesant      3714
Harlem                  2658
Bushwick               2465
Upper West Side        1971
...
Willowbrook            1
Fort Wadsworth         1
Richmondtown           1
Woodrow                1
New Dorp               1
Name: neighbourhood, Length: 221, dtype: int64
```

```
In [28]: data.neighbourhood.hist()
```

```
Out[28]: <matplotlib.axes._subplots.AxesSubplot at 0x20206d18508>
```



применим частное кодирование

```
In [29]: data['neighbourhood_freq'] = data['neighbourhood'].map(data['neighbourhood'].value_counts(normalize=True))
```

**root\_type**

```
In [30]: data.room_type.value_counts()
```

```
Out[30]: Entire home/apt    25409
Private room    22326
Shared room     1160
Name: room_type, dtype: int64
```

применим one\_hot\_encoding

```
In [31]: data = pd.get_dummies(data, columns=['room_type'], prefix='room_type', drop_f
irst=True)
```

## reviews\_per\_month

```
In [32]: data[data.reviews_per_month.isnull()].shape
```

```
Out[32]: (10052, 16)
```

заменяем null на 0 в первом случае, предположив, что null означает 0, на среднее значение, и на случайное, чтобы в последующем сравнить результаты.

```
In [33]: def impute_NA_with_random(data, NA_col=None, random_state=0):
        """Заполняем пропуски случайными значениями из этой колонки."""
        NA_col = NA_col or []
        data_copy = data.copy(deep=True)
        for i in NA_col:
            if data_copy[i].isnull().sum() > 0:
                data_copy[f'{i}_random'] = data_copy[i]
                random_sample = data_copy[i].dropna().sample(data_copy[i].isnull()
                    .sum(), random_state=random_state)
                random_sample.index = data_copy[data_copy[i].isnull()].index
                data_copy.loc[data_copy[i].isnull(), f'{i}_random'] = random_sampl
            e
        else:
            warn("Нет пропущенных значений" % i)
        return data_copy
```

```
In [34]: data = impute_NA_with_random(data=data, NA_col=['reviews_per_month'])
```

```
In [35]: data['reviews_per_month_mean'] = data['reviews_per_month'].fillna(data['review
s_per_month'].mean())
```


```
In [36]: data['reviews_per_month_zero'] = data['reviews_per_month'].fillna(0)
```

```
In [ ]:
```

## price

```
In [37]: data.boxplot(['price'])
```

```
Out[37]: <matplotlib.axes._subplots.AxesSubplot at 0x2020dc87a48>
```




переменная имеет выбросы и 0 значения, а так же распределена ненормально. предлагаю удалить значения = 0, а также прологарифмировать целевую переменную чтобы распределение стало норм - проверим это.

```
In [38]: data = data[data.price > 0]
```



```
In [39]: np.log(data['price']).hist(bins = 100)
```

```
Out[39]: <matplotlib.axes._subplots.AxesSubplot at 0x2020dc9f1c8>
```




уже похоже на нормальное распределение

```
In [40]: data['log_price'] = np.log(data['price'])
```


```
In [41]: sns.histplot(data=data, x="log_price", kde=True)
```

```
Out[41]: <matplotlib.axes._subplots.AxesSubplot at 0x2020de95d88>
```




```
In [42]: sns.distplot(data['log_price'], fit=norm)
```

```
Out[42]: <matplotlib.axes._subplots.AxesSubplot at 0x2020e38f548>
```



```
In [43]: stats.probplot(data['log_price'], plot=plt)
plt.show()
```



видим, что целевая переменная после логарифмирования распределена нормально

## Часть 3. Моделирование и улучшение качества модели с помощью future ingenering (baseline, доп фичи)

функция для записи метрик качества

```
In [44]: def mean_absolute_percentage_error(y_true, y_pred):
y_true, y_pred = np.array(y_true), np.array(y_pred)
return np.mean(np.abs((y_true - y_pred) / y_true)) * 100

def dataframe_metrics(y_test, y_pred):
    stats = [
        metrics.mean_absolute_error(y_test, y_pred),
        np.sqrt(metrics.mean_squared_error(y_test, y_pred)),
        metrics.r2_score(y_test, y_pred),
        mean_absolute_percentage_error(y_test, y_pred)
    ]
    return stats

measured_metrics = pd.DataFrame({"error_type": ["MAE", "RMSE", "R2", "MAPE"]})
measured_metrics.set_index("error_type")
```

Out[44]:

error_type
MAE
RMSE
R2
MAPE

## baseline (предсказания = median, mean)

```
In [45]: data.columns.to_list()
```

```
Out[45]: ['neighbourhood',
'latitude',
'longitude',
'price',
'minimum_nights',
'number_of_reviews',
'reviews_per_month',
'calculated_host_listings_count',
'availability_365',
'neighbourhood_group_Brooklyn',
'neighbourhood_group_Manhattan',
'neighbourhood_group_Queens',
'neighbourhood_group_Staten Island',
'neighbourhood_freq',
'room_type_Private room',
'room_type_Shared room',
'reviews_per_month_random',
'reviews_per_month_mean',
'reviews_per_month_zero',
'log_price']
```

```
In [46]: X_train, X_test, y_train, y_test = train_test_split(
        data.drop(['price', 'log_price', 'neighbourhood', 'reviews_per_month', 'review
s_per_month_random', 'reviews_per_month_mean'], axis=1),
        data['log_price'],
        test_size=0.3
    )
```

```
In [47]: X_train.shape, X_test.shape
```

```
Out[47]: ((34218, 14), (14666, 14))
```

```
In [48]: median_train = y_train.median()
        mean_train = y_train.mean()
```

```
In [49]: y_test_baseline_median = np.array([median_train]*len(y_test))
        y_test_baseline_mean = np.array([mean_train]*len(y_test))
```

```
In [50]: measured_metrics["baseline_median"] = dataframe_metrics(y_test, y_test_baselin
e_median)
        measured_metrics["baseline_mean"] = dataframe_metrics(y_test, y_test_baseline_
mean)
        measured_metrics
```

```
Out[50]:
```

	error_type	baseline_median	baseline_mean
0	MAE	0.552621	0.553657
1	RMSE	0.701938	0.697842
2	R2	-0.011807	-0.000032
3	MAPE	11.739620	11.942985

## baseline log\_reg, lassoCV, rigeCV, elasticnet\_CV

```
In [51]: continuous_vars = [
        'latitude',
        'longitude',
        'minimum_nights',
        'number_of_reviews',
        'reviews_per_month_zero',
        'calculated_host_listings_count',
        'availability_365',
    ]
```

```
In [52]: scaler = StandardScaler()
```

```
In [53]: X_train[continuous_vars] = scaler.fit_transform(X_train[continuous_vars])
        X_test[continuous_vars] = scaler.transform(X_test[continuous_vars])
```

In [54]: `X_train.head()`

Out[54]:

	latitude	longitude	minimum_nights	number_of_reviews	calculated_host_listings_count
<b>2610</b>	-0.743979	1.651368	1.041079	1.518163	-0.094473
<b>24625</b>	0.182619	-0.979271	-0.250245	0.016914	-0.186170
<b>37480</b>	1.704114	-0.037515	1.140412	-0.520847	-0.155604
<b>16803</b>	0.455341	0.439858	-0.200578	0.375421	-0.155604
<b>22886</b>	-0.619431	0.267095	-0.250245	0.532268	-0.186170

In [55]: `lin_reg = LinearRegression()  
lin_reg.fit(X_train, y_train)`

Out[55]: `LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)`


In [56]: `measured_metrics["lin_reg"] = dataframe_metrics(y_test, lin_reg.predict(X_test))  
measured_metrics`

Out[56]:

	error_type	baseline_median	baseline_mean	lin_reg
<b>0</b>	MAE	0.552621	0.553657	0.360964
<b>1</b>	RMSE	0.701938	0.697842	0.495772
<b>2</b>	R2	-0.011807	-0.000032	0.495265
<b>3</b>	MAPE	11.739620	11.942985	7.628493

```
In [57]: featureImportance = pd.DataFrame({"feature": X_train.columns,
                                           "importance": lin_reg.coef_})

featureImportance.set_index('feature', inplace=True)
featureImportance.sort_values(["importance"], ascending=False, inplace=True)
featureImportance["importance"].plot(kind = 'bar', figsize=(20, 6));
```




## Ridge CV

```
In [58]: ridge_cv = RidgeCV()
ridge_cv.fit(X_train, y_train)
```

```
Out[58]: RidgeCV(alphas=array([ 0.1,  1. , 10. ]), cv=None, fit_intercept=True,
                  gcvcv_mode=None, normalize=False, scoring=None, store_cv_values=False)
```

```
In [59]: featureImportance = pd.DataFrame({"feature": X_train.columns,
                                           "importance": ridge_cv.coef_})

featureImportance.set_index('feature', inplace=True)
featureImportance.sort_values(["importance"], ascending=False, inplace=True)
featureImportance["importance"].plot(kind = 'bar', figsize=(18, 6));
```



```
In [60]: measured_metrics["ridge_cv"] = dataframe_metrics(y_test, ridge_cv.predict(X_test))
measured_metrics
```

Out[60]:

	error_type	baseline_median	baseline_mean	lin_reg	ridge_cv
0	MAE	0.552621	0.553657	0.360964	0.360964
1	RMSE	0.701938	0.697842	0.495772	0.495773
2	R2	-0.011807	-0.000032	0.495265	0.495262
3	MAPE	11.739620	11.942985	7.628493	7.628503

## Lasso CV


```
In [61]: lasso_cv = LassoCV()
lasso_cv.fit(X_train, y_train)
```

```
Out[61]: LassoCV(alphas=None, copy_X=True, cv='warn', eps=0.001, fit_intercept=True,
max_iter=1000, n_alphas=100, n_jobs=None, normalize=False,
positive=False, precompute='auto', random_state=None,
selection='cyclic', tol=0.0001, verbose=False)
```



```
In [62]: featureImportance = pd.DataFrame({"feature": X_train.columns[lasso_cv.coef_!=0
],
                                           "importance": lasso_cv.coef_[lasso_cv.coef_!
=0]})

featureImportance.set_index('feature', inplace=True)
featureImportance.sort_values(["importance"], ascending=False, inplace=True)
featureImportance["importance"].plot(kind = 'bar', figsize=(18, 6));
```



```
In [63]: measured_metrics["lasso_cv"] = dataframe_metrics(y_test, lasso_cv.predict(X_te
st))
measured_metrics
```

Out[63]:

	error_type	baseline_median	baseline_mean	lin_reg	ridge_cv	lasso_cv
0	MAE	0.552621	0.553657	0.360964	0.360964	0.361060
1	RMSE	0.701938	0.697842	0.495772	0.495773	0.495993
2	R2	-0.011807	-0.000032	0.495265	0.495262	0.494815
3	MAPE	11.739620	11.942985	7.628493	7.628503	7.631662


## ElasticNet CV

```
In [64]: elastic_cv = ElasticNetCV()
elastic_cv.fit(X_train, y_train)
```

```
Out[64]: ElasticNetCV(alphas=None, copy_X=True, cv='warn', eps=0.001, fit_intercept=True,
                        l1_ratio=0.5, max_iter=1000, n_alphas=100, n_jobs=None,
                        normalize=False, positive=False, precompute='auto',
                        random_state=None, selection='cyclic', tol=0.0001, verbose=0)
```

```
In [65]: featureImportance = pd.DataFrame({"feature": X_train.columns[elastic_cv.coef_!=0],
                                           "importance": elastic_cv.coef_[elastic_cv.coef_!=0]})

featureImportance.set_index('feature', inplace=True)
featureImportance.sort_values(["importance"], ascending=False, inplace=True)
featureImportance["importance"].plot(kind='bar', figsize=(18, 6));
```



```
In [66]: measured_metrics["elastic_cv"] = dataframe_metrics(y_test, elastic_cv.predict(
X_test))
measured_metrics
```

Out[66]:

	error_type	baseline_median	baseline_mean	lin_reg	ridge_cv	lasso_cv	elastic_cv
0	MAE	0.552621	0.553657	0.360964	0.360964	0.361060	0.361257
1	RMSE	0.701938	0.697842	0.495772	0.495773	0.495993	0.496253
2	R2	-0.011807	-0.000032	0.495265	0.495262	0.494815	0.494284
3	MAPE	11.739620	11.942985	7.628493	7.628503	7.631662	7.636934

## улучшение модели (обработка latitude, longitude)

возьмем за центр нью-йорка - центр Central Park (40.782748, -73.965743), создадим новую фичу - расстояние от недвижимости до центра Central Park

```
In [67]: data['dist_manh'] = np.sqrt((data['latitude'] - 40.782748)**2 + (data['longitude'] - (-73.965743))**2)
```

In [68]: data.head()

Out[68]:

	neighbourhood	latitude	longitude	price	minimum_nights	number_of_reviews	reviews_per_
0	Kensington	40.64749	-73.97237	149	1	9	
1	Midtown	40.75362	-73.98377	225	1	45	
2	Harlem	40.80902	-73.94190	150	3	0	
3	Clinton Hill	40.68514	-73.95976	89	1	270	
4	East Harlem	40.79851	-73.94399	80	10	9	

```

In [69]: X_train, X_test, y_train, y_test = train_test_split(
    data.drop(['price', 'log_price', 'neighbourhood', 'latitude', 'longitude', 'reviews_per_month', 'reviews_per_month_random', 'reviews_per_month_mean'], axis=1),
    data['log_price'],
    test_size=0.3
)
continuous_vars = [
    'dist_manh',
    'minimum_nights',
    'number_of_reviews',
    'reviews_per_month_zero',
    'calculated_host_listings_count',
    'availability_365',
]
scaler = StandardScaler()
X_train[continuous_vars] = scaler.fit_transform(X_train[continuous_vars])
X_test[continuous_vars] = scaler.transform(X_test[continuous_vars])
lin_reg = LinearRegression()
lin_reg.fit(X_train, y_train)
ridge_cv = RidgeCV()
ridge_cv.fit(X_train, y_train)
lasso_cv = LassoCV()
lasso_cv.fit(X_train, y_train)
elastic_cv = ElasticNetCV()
elastic_cv.fit(X_train, y_train)
measured_metrics["lin_reg_manh_dist"] = dataframe_metrics(y_test, lin_reg.predict(X_test))
measured_metrics["ridge_cv_manh_dist"] = dataframe_metrics(y_test, ridge_cv.predict(X_test))
measured_metrics["lasso_cv_manh_dist"] = dataframe_metrics(y_test, lasso_cv.predict(X_test))
measured_metrics

```

Out[69]:

	error_type	baseline_median	baseline_mean	lin_reg	ridge_cv	lasso_cv	elastic_cv	lin_reg
0	MAE	0.552621	0.553657	0.360964	0.360964	0.361060	0.361257	
1	RMSE	0.701938	0.697842	0.495772	0.495773	0.495993	0.496253	
2	R2	-0.011807	-0.000032	0.495265	0.495262	0.494815	0.494284	
3	MAPE	11.739620	11.942985	7.628493	7.628503	7.631662	7.636934	

новая фича не улучшила модель

**улучшение модели (обработка min\_max\_scaller)**

```

In [70]: X_train, X_test, y_train, y_test = train_test_split(
    data.drop(['price', 'log_price', 'neighbourhood', 'dist_manh', 'reviews_per_mo
nth', 'reviews_per_month_random', 'reviews_per_month_mean'], axis=1),
    data['log_price'],
    test_size=0.3
)
continuous_vars = [
    'latitude',
    'longitude',
    'minimum_nights',
    'number_of_reviews',
    'reviews_per_month_zero',
    'calculated_host_listings_count',
    'availability_365',
]
scaler = MinMaxScaler()
X_train[continuous_vars] = scaler.fit_transform(X_train[continuous_vars])
X_test[continuous_vars] = scaler.transform(X_test[continuous_vars])
lin_reg = LinearRegression()
lin_reg.fit(X_train, y_train)
ridge_cv = RidgeCV()
ridge_cv.fit(X_train, y_train)
lasso_cv = LassoCV()
lasso_cv.fit(X_train, y_train)
elastic_cv = ElasticNetCV()
elastic_cv.fit(X_train, y_train)
measured_metrics["lin_reg_minmax"] = dataframe_metrics(y_test, lin_reg.predict
(X_test))
measured_metrics["ridge_cv_minmax"] = dataframe_metrics(y_test, ridge_cv.predi
ct(X_test))
measured_metrics["lasso_cv_minmax"] = dataframe_metrics(y_test, lasso_cv.predi
ct(X_test))
measured_metrics

```

Out[70]:

	error_type	baseline_median	baseline_mean	lin_reg	ridge_cv	lasso_cv	elastic_cv	lin_reg
0	MAE	0.552621	0.553657	0.360964	0.360964	0.361060	0.361257	
1	RMSE	0.701938	0.697842	0.495772	0.495773	0.495993	0.496253	
2	R2	-0.011807	-0.000032	0.495265	0.495262	0.494815	0.494284	
3	MAPE	11.739620	11.942985	7.628493	7.628503	7.631662	7.636934	

## улучшение модели robustscaller

```

In [71]: X_train, X_test, y_train, y_test = train_test_split(
    data.drop(['price', 'log_price', 'neighbourhood', 'dist_manh', 'reviews_per_mo
nth', 'reviews_per_month_random', 'reviews_per_month_mean'], axis=1),
    data['log_price'],
    test_size=0.3
)
continuous_vars = [
    'latitude',
    'longitude',
    'minimum_nights',
    'number_of_reviews',
    'reviews_per_month_zero',
    'calculated_host_listings_count',
    'availability_365',
]
scaler = RobustScaler()
X_train[continuous_vars] = scaler.fit_transform(X_train[continuous_vars])
X_test[continuous_vars] = scaler.transform(X_test[continuous_vars])
lin_reg = LinearRegression()
lin_reg.fit(X_train, y_train)
ridge_cv = RidgeCV()
ridge_cv.fit(X_train, y_train)
lasso_cv = LassoCV()
lasso_cv.fit(X_train, y_train)
elastic_cv = ElasticNetCV()
elastic_cv.fit(X_train, y_train)
measured_metrics["lin_reg_robust"] = dataframe_metrics(y_test, lin_reg.predict
(X_test))
measured_metrics["ridge_cv_robust"] = dataframe_metrics(y_test, ridge_cv.predi
ct(X_test))
measured_metrics["lasso_cv_robust"] = dataframe_metrics(y_test, lasso_cv.predi
ct(X_test))
measured_metrics

```

Out[71]:

	error_type	baseline_median	baseline_mean	lin_reg	ridge_cv	lasso_cv	elastic_cv	lin_reg
0	MAE	0.552621	0.553657	0.360964	0.360964	0.361060	0.361257	
1	RMSE	0.701938	0.697842	0.495772	0.495773	0.495993	0.496253	
2	R2	-0.011807	-0.000032	0.495265	0.495262	0.494815	0.494284	
3	MAPE	11.739620	11.942985	7.628493	7.628503	7.631662	7.636934	

**улучшение модели (убираем фичу available\_365)**

```

In [72]: X_train, X_test, y_train, y_test = train_test_split(
    data.drop(['availability_365', 'price', 'log_price', 'neighbourhood', 'dist_ma
nh', 'reviews_per_month', 'reviews_per_month_random', 'reviews_per_month_mean'],
    axis=1),
    data['log_price'],
    test_size=0.3
)
continuous_vars = [
    'latitude',
    'longitude',
    'minimum_nights',
    'number_of_reviews',
    'reviews_per_month_zero',
    'calculated_host_listings_count',
]
scaler = RobustScaler()
X_train[continuous_vars] = scaler.fit_transform(X_train[continuous_vars])
X_test[continuous_vars] = scaler.transform(X_test[continuous_vars])
lin_reg = LinearRegression()
lin_reg.fit(X_train, y_train)
ridge_cv = RidgeCV()
ridge_cv.fit(X_train, y_train)
lasso_cv = LassoCV()
lasso_cv.fit(X_train, y_train)
elastic_cv = ElasticNetCV()
elastic_cv.fit(X_train, y_train)
measured_metrics["lin_reg_no_365"] = dataframe_metrics(y_test, lin_reg.predict
(X_test))
measured_metrics["ridge_cv_no_365"] = dataframe_metrics(y_test, ridge_cv.predi
ct(X_test))
measured_metrics["lasso_cv_no_365"] = dataframe_metrics(y_test, lasso_cv.predi
ct(X_test))
measured_metrics

```

Out[72]:

	error_type	baseline_median	baseline_mean	lin_reg	ridge_cv	lasso_cv	elastic_cv	lin_reg
0	MAE	0.552621	0.553657	0.360964	0.360964	0.361060	0.361257	
1	RMSE	0.701938	0.697842	0.495772	0.495773	0.495993	0.496253	
2	R2	-0.011807	-0.000032	0.495265	0.495262	0.494815	0.494284	
3	MAPE	11.739620	11.942985	7.628493	7.628503	7.631662	7.636934	

## пробы random\_forest

```

In [73]: from sklearn.ensemble import RandomForestRegressor

```

```
In [74]: rf_reg = RandomForestRegressor(n_jobs=-1)
rf_reg.fit(X_train, y_train)

measured_metrics["rf_reg"] = dataframe_metrics(y_test, rf_reg.predict(X_test))
measured_metrics
```

Out[74]:

cv_minmax	lin_reg_robust	ridge_cv_robust	lasso_cv_robust	lin_reg_no_365	ridge_cv_no_365	lasso_cv_no_365
0.362994	0.360272	0.360274	0.363005	0.368039	0.368037	0.368037
0.499830	0.495004	0.495002	0.498898	0.510643	0.510640	0.510640
0.483971	0.495907	0.495912	0.487944	0.472819	0.472826	0.472826
7.677352	7.594587	7.594647	7.667996	7.689093	7.689060	7.689060

In [ ]: