

▼ Чем лучше бустить? Тестируем алгоритмы бустинга в бою.

```
from google.colab import drive
drive.mount('/content/drive')
```

```
Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).
```

```
!pip install catboost
```

```
Requirement already satisfied: catboost in /usr/local/lib/python3.7/dist-packages (0.25)
Requirement already satisfied: numpy>=1.16.0 in /usr/local/lib/python3.7/dist-packages (from catboost) (1.19.5)
Requirement already satisfied: graphviz in /usr/local/lib/python3.7/dist-packages (from catboost) (0.10.1)
Requirement already satisfied: scipy in /usr/local/lib/python3.7/dist-packages (from catboost) (1.4.1)
Requirement already satisfied: six in /usr/local/lib/python3.7/dist-packages (from catboost) (1.15.0)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.7/dist-packages (from catboost) (3.2.2)
Requirement already satisfied: pandas>=0.24.0 in /usr/local/lib/python3.7/dist-packages (from catboost) (1.1.5)
Requirement already satisfied: plotly in /usr/local/lib/python3.7/dist-packages (from catboost) (4.4.1)
Requirement already satisfied: cyclers>=0.10 in /usr/local/lib/python3.7/dist-packages (from matplotlib->catboost) (0.10.0)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib->catbo
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib->catboost) (2.8.1)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib->catboost) (1.3.1)
Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.7/dist-packages (from pandas>=0.24.0->catboost) (2018.9)
Requirement already satisfied: retrying>=1.3.3 in /usr/local/lib/python3.7/dist-packages (from plotly->catboost) (1.3.3)
```

Часть 1. EDA, Часть 2. Preprocessing & Feature Engineering

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
```

```
from sklearn.preprocessing import LabelEncoder
```

```

from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split,cross_val_score, StratifiedKFold, GridSearchCV
from sklearn.metrics import accuracy_score,confusion_matrix,roc_auc_score,roc_curve,classification_report, precision_score, recall_score, f1_score
from xgboost.sklearn import XGBClassifier
from catboost import CatBoostClassifier
from lightgbm import LGBMClassifier

```

```

%matplotlib inline
plt.rcParams["figure.figsize"] = (12,8)

```

```

data = pd.read_csv('/content/drive/MyDrive/STUDY/otus/HW/3/WA_Fn-UseC_-Telco-Customer-Churn.csv')

```

```

data.head()

```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL	No	No
1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL	Yes	No
2	3668-QPYBK	Male	0	No	No	2	Yes	No	DSL	Yes	No
3	7795-CFOCW	Male	0	No	No	45	No	No phone service	DSL	Yes	No
4	9237-HQITU	Female	0	No	No	2	Yes	No	Fiber optic	No	No

проверка на дубликаты

```

data[data.duplicated(['customerID'], keep=False)]

```

customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBack
------------	--------	---------------	---------	------------	--------	--------------	---------------	-----------------	----------------	------------

data.describe()

	SeniorCitizen	tenure	MonthlyCharges
count	7043.000000	7043.000000	7043.000000
mean	0.162147	32.371149	64.761692
std	0.368612	24.559481	30.090047
min	0.000000	0.000000	18.250000
25%	0.000000	9.000000	35.500000
50%	0.000000	29.000000	70.350000
75%	0.000000	55.000000	89.850000
max	1.000000	72.000000	118.750000

data.columns

```
Index(['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents',
      'tenure', 'PhoneService', 'MultipleLines', 'InternetService',
      'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport',
      'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBilling',
      'PaymentMethod', 'MonthlyCharges', 'TotalCharges', 'Churn'],
      dtype='object')
```

data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column              Non-Null Count  Dtype
---  -
0   customerID          7043 non-null  object
1   gender              7043 non-null  object
```

```
2 SeniorCitizen      7043 non-null  int64
3 Partner            7043 non-null  object
4 Dependents         7043 non-null  object
5 tenure             7043 non-null  int64
6 PhoneService       7043 non-null  object
7 MultipleLines      7043 non-null  object
8 InternetService    7043 non-null  object
9 OnlineSecurity     7043 non-null  object
10 OnlineBackup       7043 non-null  object
11 DeviceProtection  7043 non-null  object
12 TechSupport       7043 non-null  object
13 StreamingTV       7043 non-null  object
14 StreamingMovies   7043 non-null  object
15 Contract          7043 non-null  object
16 PaperlessBilling  7043 non-null  object
17 PaymentMethod     7043 non-null  object
18 MonthlyCharges    7043 non-null  float64
19 TotalCharges      7043 non-null  object
20 Churn             7043 non-null  object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

data.tail()

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	Online
7038	6840-RESVB	Male	0	Yes	Yes	24	Yes	Yes	DSL	Yes	
7039	2234-XADUH	Female	0	Yes	Yes	72	Yes	Yes	Fiber optic	No	
7040	4801-JZAZL	Female	0	Yes	Yes	11	No	No phone service	DSL	Yes	
7041	8361-LTMKD	Male	1	Yes	No	4	Yes	Yes	Fiber optic	No	
7042	3186-AJIEK	Male	0	No	No	66	Yes	No	Fiber optic	Yes	

проверим на пропуски

```
data.isnull().sum()
```




```
customerID      0
gender          0
SeniorCitizen   0
Partner         0
Dependents      0
tenure          0
PhoneService    0
MultipleLines   0
InternetService 0
OnlineSecurity  0
OnlineBackup    0
DeviceProtection 0
TechSupport     0
StreamingTV     0
StreamingMovies 0
Contract        0
PaperlessBilling 0
PaymentMethod   0
MonthlyCharges  0
TotalCharges    0
Churn           0
dtype: int64
```

удалим лишнюю колонку

```
data.drop(columns=['customerID'], inplace = True)
```


```
data.hist()
```

```
array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7f4a20ff90d0>,  
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f4a20fc8710>],  
      [<matplotlib.axes._subplots.AxesSubplot object at 0x7f4a20f82d90>,  
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f4a20f45450>]],  
      dtype=object)
```



```
sns.pairplot(data)
```

<seaborn.axisgrid.PairGrid at 0x7f4a20dd05d0>



Target: Churn

fig = sns.pairplot(data, hue='Churn', palette='magma')

```
data['Churn'].value_counts().plot(kind='bar', label='Churn').legend()
plt.title('Распределение оттока клиентов')
```

```
Text(0.5, 1.0, 'Распределение оттока клиентов')
```



оцифруем целевую переменную


```
data['Churn'] = data['Churn'].apply(lambda x: 1 if x == 'Yes' else 0)
```

▼ gender

```
data.gender.value_counts()
```

```
Male      3555
Female    3488
Name: gender, dtype: int64
```

```
sns.countplot(x='gender', hue='Churn', data=data);
```

закодируем, используя label encoder. Его т.к. будем работать с "деревянными" моделями


```
le = LabelEncoder()
data.gender = le.fit_transform(data.gender)
```

SeniorCitizen

```
data.SeniorCitizen.value_counts()

0    5901
1    1142
Name: SeniorCitizen, dtype: int64
```

```
sns.countplot(x='SeniorCitizen', hue='Churn', data=data);
```




```
data.SeniorCitizen = le.fit_transform(data.SeniorCitizen)
```

▼ Partner

```
data.Partner.value_counts()
```

```
No      3641
Yes     3402
Name: Partner, dtype: int64
```

```
sns.countplot(x='Partner', hue='Churn', data=data);
```



интересный признак, люди без партнера чаще уходят от оператора, возможно это связано с тем, что одному проще сменить оператора. В паре нужно менять всем, т.к. звонки между одним и тем же оператором - дешевле


```
data.Partner = le.fit_transform(data.Partner)
```

Dependents

```
data.Dependents.value_counts()
```

```
No      4933  
Yes      2110  
Name: Dependents, dtype: int64
```

```
sns.countplot(x='Dependents', hue='Churn', data=data);
```




клиенты без иждивенцев чаще уходят

```
data.Dependents = le.fit_transform(data.Dependents)
```

▼ tenure


```
data.tenure.hist(bins = 100)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f4a1579ea10>
```



```
sns.boxplot(x='Churn', y='tenure', data=data)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f4a15676a50>
```




видим, что иим что уходящие пользуются связью не так долго как постоянные клиенты

▼ PhoneService

```
data.PhoneService.value_counts()
```

```
Yes    6361
No      682
Name: PhoneService, dtype: int64
```

```
sns.countplot(x='PhoneService', hue='Churn', data=data);
```




```
data.PhoneService = le.fit_transform(data.PhoneService)
```

MultipleLines

```
data.MultipleLines.value_counts()
```

```
No          3390
Yes          2971
No phone service    682
Name: MultipleLines, dtype: int64
```

```
sns.countplot(x='MultipleLines', hue='Churn', data=data):
```



```
data.MultipleLines = le.fit_transform(data.MultipleLines)
```

▼ InternetService


```
data.InternetService.value_counts()
```

```
Fiber optic    3096
DSL            2421
```



```
No          1526  
Name: InternetService, dtype: int64
```

```
sns.countplot(x='InternetService', hue='Churn', data=data);
```



интересное замечание, клиенты с оптоволоконном лидируют по оттоку..


```
data.InternetService = le.fit_transform(data.InternetService)
```

▼ OnlineSecurity

```
data.OnlineSecurity.value_counts()
```

```
No          3498  
Yes         2019  
No internet service  1526  
Name: OnlineSecurity, dtype: int64
```

```
sns.countplot(x='OnlineSecurity', hue='Churn', data=data);
```




```
data.OnlineSecurity = le.fit_transform(data.OnlineSecurity)
```

▼ OnlineBackup

```
data.OnlineBackup.value_counts()
```

```
No          3088  
Yes         2429  
No internet service  1526  
Name: OnlineBackup, dtype: int64
```

```
sns.countplot(x='OnlineBackup', hue='Churn', data=data);
```




```
data.OnlineBackup = le.fit_transform(data.OnlineBackup)
```

▼ DeviceProtection

```
data.DeviceProtection.value_counts()
```

```
No          3095  
Yes         2422  
No internet service  1526  
Name: DeviceProtection, dtype: int64
```

```
sns.countplot(x='DeviceProtection', hue='Churn', data=data);
```



data.DeviceProtection = le.fit_transform(data.DeviceProtection)



TechSupport



data.TechSupport.value_counts()

No 3473
Yes 2044
No internet service 1526
Name: TechSupport, dtype: int64



sns.countplot(x='TechSupport', hue='Churn', data=data);



```
data.TechSupport = le.fit_transform(data.TechSupport)
```



▼ StreamingTV



```
data.StreamingTV.value_counts()
```

```
No 2810
```


```
Yes 2707
```

```
No internet service 1526
```

```
Name: StreamingTV, dtype: int64
```



```
sns.countplot(x='StreamingTV', hue='Churn', data=data);
```



```
data.StreamingTV = le.fit_transform(data.StreamingTV)
```



StreamingMovies




```
data.StreamingMovies.value_counts()
```

```
No          2785
Yes          2732
No internet service  1526
Name: StreamingMovies, dtype: int64
```



```
sns.countplot(x='StreamingMovies', hue='Churn', data=data);
```



```
data.StreamingMovies = le.fit_transform(data.StreamingMovies)
```

```
data.head()
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	Device
0	0	0	1	0	1	0	1	0	0	2	
1	1	0	0	0	34	1	0	0	2	0	
2	1	0	0	0	2	1	0	0	2	2	
3	1	0	0	0	45	0	1	0	2	0	
4	0	0	0	0	2	1	0	1	0	0	


Contract

```
data.Contract.value_counts()
```

Month-to-month	3875
Two year	1695


```
One year      1473  
Name: Contract, dtype: int64
```

```
sns.countplot(x='Contract', hue='Churn', data=data);
```



крутой признак, месячники уходят чаще


```
data.Contract = le.fit_transform(data.Contract)
```

▼ PaperlessBilling

```
data.PaperlessBilling.value_counts()
```

```
Yes      4171  
No       2872  
Name: PaperlessBilling, dtype: int64
```

```
sns.countplot(x='PaperlessBilling', hue='Churn', data=data);
```



```
data.PaperlessBilling = le.fit_transform(data.PaperlessBilling)
```

▼ PaymentMethod

```
data.PaymentMethod.value_counts()
```

```
Electronic check      2365  
Mailed check          1612  
Bank transfer (automatic)  1544  
Credit card (automatic)  1522  
Name: PaymentMethod, dtype: int64
```

```
sns.countplot(x='PaymentMethod', hue='Churn', data=data);
```

электронные чеки это плохо...



```
data.PaymentMethod = le.fit_transform(data.PaymentMethod)
```



```
data.head()
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	Device
0	0	0	1	0	1	0	1	0	0	2	
1	1	0	0	0	34	1	0	0	2	0	
2	1	0	0	0	2	1	0	0	2	2	
3	1	0	0	0	45	0	1	0	2	0	
4	0	0	0	0	2	1	0	1	0	0	





MonthlyCharges TotalCharges



```
_, axes = plt.subplots(1, 2, sharey=True, figsize=(16,6))

sns.boxplot(x='Churn', y='MonthlyCharges', data=data, ax=axes[0]);
sns.violinplot(x='Churn', y='MonthlyCharges', data=data, ax=axes[1]);
```



преобразуем total_charges во float

```
data['TotalCharges']
```

```
0      29.85
1    1889.5
2     108.15
3   1840.75
4     151.65
...
7038   1990.5
7039   7362.9
7040    346.45
7041    306.6
7042   6844.5
```

Name: TotalCharges, Length: 7043, dtype: object

```
data[data['TotalCharges'].str.match(' ') == False]
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	Dev
0	0	0	1	0	1	0	1	0	0	2	
1	1	0	0	0	34	1	0	0	2	0	
2	1	0	0	0	2	1	0	0	2	2	
3	1	0	0	0	45	0	1	0	2	0	
4	0	0	0	0	2	1	0	1	0	0	
...	
7038	1	0	1	1	24	1	2	0	2	0	
7039	0	0	1	1	72	1	2	1	0	2	

```
data = data[data['TotalCharges'].str.match(' ') == False]
```

```
data['TotalCharges'] = data['TotalCharges'].astype('float64')
```


```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
"""Entry point for launching an IPython kernel.
```

```
_, axes = plt.subplots(1, 2, sharey=True, figsize=(16,6))
```


```
sns.boxplot(x='Churn', y='TotalCharges', data=data, ax=axes[0]);
sns.violinplot(x='Churn', y='TotalCharges', data=data, ax=axes[1]);
```




```
data.TotalCharges.corr(data.MonthlyCharges)
```

```
0.6510648032262024
```

```
data.MonthlyCharges.hist(bins = 100)
```




```
<matplotlib.axes._subplots.AxesSubplot at 0x7f4a152d2f10>
```



```
data.TotalCharges.hist(bins = 100)
```


```
<matplotlib.axes._subplots.AxesSubplot at 0x7f4a1318de10>
```



```
corr_matrix = data.corr()
```



```
sns.heatmap(corr_matrix);
```



Часть 3. Who's the mightiest of them all?

функция качества модели

```
def quality(prediction_y, true_y):
    accuracy = accuracy_score(prediction_y, true_y)
    precision = precision_score(prediction_y, true_y)
    recall = recall_score(prediction_y, true_y)
    f1 = f1_score(prediction_y, true_y)
    print("Accuracy: {:.3f}\nPrecision: {:.3f}\nRecall: {:.3f}\nF1-score: {:.3f}".format(
        accuracy, precision, recall, f1
    ))
```

функция построения кривой roc_auc

```
def plot_roc_curve(prob_prediction, actual):
    fpr, tpr, thresholds = roc_curve(y_test, prob_prediction)
    auc_score = roc_auc_score(y_test, prob_prediction)

    plt.plot(fpr, tpr, label='ROC curve ')
    plt.plot([0, 1], [0, 1])
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('ROC AUC: {:.3f}'.format(auc_score))
    plt.show()
```

▼ sklearn RandomForest

```
X_train, X_test, y_train, y_test = train_test_split(
    data.drop(['Churn'], axis = 1), data.Churn, test_size=0.3, random_state=2021, stratify=data.Churn.values)
```

```
clfRF = RandomForestClassifier()
clfRF.fit(X_train, y_train)
```

```
RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                        criterion='gini', max_depth=None, max_features='auto',
```

```
max_leaf_nodes=None, max_samples=None,  
min_impurity_decrease=0.0, min_impurity_split=None,  
min_samples_leaf=1, min_samples_split=2,  
min_weight_fraction_leaf=0.0, n_estimators=100,  
n_jobs=None, oob_score=False, random_state=None,  
verbose=0, warm_start=False)
```


```
predRF = clfRF.predict(X_test)
```

```
print("Train quality:")  
quality(clfRF.predict(X_train), y_train)  
print("\nTest quality:")  
quality(predRF, y_test)
```

```
Train quality:  
Accuracy:  0.998  
Precision: 0.994  
Recall:    0.997  
F1-score:  0.995
```

```
Test quality:  
Accuracy:  0.779  
Precision: 0.453  
Recall:    0.614  
F1-score:  0.521
```

```
plot_roc_curve(predRF, y_test)
```



▼ XGBoost


```
clfXGB = XGBClassifier()  
clfXGB.fit(X_train, y_train)  
predXGB = clfXGB.predict(X_test)
```

```
print("Train quality:")  
quality(clfXGB.predict(X_train), y_train)  
print("\nTest quality:")  
quality(predXGB, y_test)
```

```
Train quality:  
Accuracy:  0.825  
Precision: 0.563  
Recall:    0.719  
F1-score:  0.631
```

```
Test quality:  
Accuracy:  0.799  
Precision: 0.512  
Recall:    0.657  
F1-score:  0.575
```

```
plot_roc_curve(predXGB, y_test)
```



▼ CatBoost

```
clfcat = CatBoostClassifier(eval_metric='AUC')  
clfcat.fit(X_train, y_train, silent= True)  
predcat = clfcat.predict(X_test)
```

```
print("Train quality:")
```

```
print('Train quality: ',  
      quality(clfcat.predict(X_train), y_train)  
      )  
print("\nTest quality:")  
quality(predcat, y_test)
```

```
Train quality:  
Accuracy:  0.880  
Precision: 0.694  
Recall:    0.828  
F1-score:  0.755
```

```
Test quality:  
Accuracy:  0.789  
Precision: 0.506  
Recall:    0.627  
F1-score:  0.560
```

```
plot_roc_curve(predcat, y_test)
```

ROC AUC: 0.699



▼ LightGBM


```
clfLBM = LGBMClassifier()  
clfLBM.fit(X_train, y_train)  
predLBM = clfLBM.predict(X_test)
```

```
print("Train quality:")  
quality(clfLBM.predict(X_train), y_train)  
print("\nTest quality:")  
quality(predLBM, y_test)
```

```
Train quality:  
Accuracy:  0.889  
Precision: 0.722  
Recall:    0.836  
F1-score:  0.775
```

```
Test quality:  
Accuracy:  0.784  
Precision: 0.497  
Recall:    0.616  
F1-score:  0.550
```

```
plot_roc_curve(predLBM, y_test)
```



sklearn RandomForest, XGBoost, CatBost, LightGBM (настройка гиперпараметров)



Инициализируем страифицированную разбивку нашего датасета для валидации

```
skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
```

sklearn RandomForest

```
parameters = {'max_features': [4, 7, 10, 13], 'min_samples_leaf': [1, 3, 5, 7], 'max_depth': [5,10,15,20]}
rfc = RandomForestClassifier(n_estimators=100, random_state=42,
                             n_jobs=-1, oob_score=True)
gcv_rf = GridSearchCV(rfc, parameters, n_jobs=-1, cv=skf, verbose=1)
gcv_rf.fit(X_train, y_train)
```



```

Fitting 5 folds for each of 64 candidates, totalling 320 fits
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done 46 tasks      | elapsed: 19.2s
[Parallel(n_jobs=-1)]: Done 196 tasks    | elapsed: 1.6min
[Parallel(n_jobs=-1)]: Done 320 out of 320 | elapsed: 2.9min finished
GridSearchCV(cv=StratifiedKFold(n_splits=5, random_state=42, shuffle=True),
             error_score=nan,
             estimator=RandomForestClassifier(bootstrap=True, ccp_alpha=0.0,
                                                class_weight=None,
                                                criterion='gini', max_depth=None,
                                                max_features='auto',
                                                max_leaf_nodes=None,
                                                max_samples=None,
                                                min_impurity_decrease=0.0,
                                                min_impurity_split=None,
                                                min_samples_leaf=1,
                                                min_samples_split=2,
                                                min_weight_fraction_leaf=0.0,
                                                n_estimators=100, n_jobs=-1,
                                                oob_score=True, random_state=42,
                                                verbose=0, warm_start=False),
             iid='deprecated', n_jobs=-1,
             param_grid={'max_depth': [5, 10, 15, 20],
                          'max_features': [4, 7, 10, 13],
                          'min_samples_leaf': [1, 3, 5, 7]},
             pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
             scoring=None, verbose=1)

```

```

print("Train quality:")
quality(gcv_rf.predict(X_train), y_train)
print("\nTest quality:")
quality(gcv_rf.predict(X_test), y_test)

```

```

Train quality:
Accuracy: 0.869
Precision: 0.633
Recall: 0.834
F1-score: 0.720

```


```

Test quality:
Accuracy: 0.791
Precision: 0.481

```

Recall: 0.643
F1-score: 0.550

```
plot_roc_curve(gcv_rf.predict(X_test), y_test)
```



XGBoost

```
xgbc = XGBClassifier(n_estimators=100, random_state=42,  
                    n_jobs=-1, oob_score=True)  
gcv_x = GridSearchCV(xgbc, parameters, n_jobs=-1, cv=skf, verbose=1)  
gcv_x.fit(X_train, y_train)
```

```

Fitting 5 folds for each of 64 candidates, totalling 320 fits
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done 46 tasks      | elapsed: 17.0s
[Parallel(n_jobs=-1)]: Done 196 tasks    | elapsed: 2.2min
[Parallel(n_jobs=-1)]: Done 320 out of 320 | elapsed: 4.9min finished
GridSearchCV(cv=StratifiedKFold(n_splits=5, random_state=42, shuffle=True),
             error_score=nan,
             estimator=XGBClassifier(base_score=0.5, booster='gbtree',
                                     colsample_bylevel=1, colsample_bynode=1,
                                     colsample_bytree=1, gamma=0,
                                     learning_rate=0.1, max_delta_step=0,
                                     max_depth=3, min_child_weight=1,
                                     missing=None, n_estimators=100, n_jobs=-1,
                                     nthread=None, objective='binary:logistic',
                                     oob_score=True, random_state=42,
                                     reg_alpha=0, reg_lambda=1,
                                     scale_pos_weight=1, seed=None, silent=None,
                                     subsample=1, verbosity=1),
             iid='deprecated', n_jobs=-1,
             param_grid={'max_depth': [5, 10, 15, 20],
                         'max_features': [4, 7, 10, 13],
                         'min_samples_leaf': [1, 3, 5, 7]},
             pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
             scoring=None, verbose=1)

```

```

print("Train quality:")
quality(gcv_x.predict(X_train), y_train)
print("\nTest quality:")
quality(gcv_x.predict(X_test), y_test)

```

```

Train quality:
Accuracy: 0.856
Precision: 0.635
Recall: 0.782
F1-score: 0.700


```

```

Test quality:
Accuracy: 0.798
Precision: 0.515
Recall: 0.652
F1-score: 0.576

```

```
plot_roc_curve(gcv_x.predict(X_test), y_test)
```



CatBoost

```
param_cat = {'iterations': [1,5,10,20,50], 'subsample': [0.66, 0.8,1], 'max_depth': [5,10,15,20]}
catc = CatBoostClassifier(random_state=42,
                           eval_metric='AUC')
gcv_c = GridSearchCV(catc, param_cat, n_jobs=-1, cv=skf, verbose=1)
gcv_c.fit(X_train, y_train)

1:      total: 3.78ms   remaining: 90.7ms
```

2:	total: 5.03ms	remaining: 78.8ms
3:	total: 6.82ms	remaining: 78.5ms
4:	total: 8.62ms	remaining: 77.6ms
5:	total: 10.4ms	remaining: 76.3ms
6:	total: 12.3ms	remaining: 75.3ms
7:	total: 14ms	remaining: 73.6ms
8:	total: 16ms	remaining: 72.7ms
9:	total: 17.8ms	remaining: 71.2ms
10:	total: 19.6ms	remaining: 69.4ms
11:	total: 21.5ms	remaining: 68ms
12:	total: 23.4ms	remaining: 66.5ms
13:	total: 25.2ms	remaining: 64.7ms
14:	total: 27ms	remaining: 63ms
15:	total: 28.8ms	remaining: 61.2ms
16:	total: 30.6ms	remaining: 59.4ms
17:	total: 32.4ms	remaining: 57.6ms
18:	total: 34.2ms	remaining: 55.8ms
19:	total: 36ms	remaining: 54ms
20:	total: 37.9ms	remaining: 52.4ms
21:	total: 39.8ms	remaining: 50.6ms
22:	total: 41.5ms	remaining: 48.8ms
23:	total: 43.3ms	remaining: 47ms
24:	total: 45.1ms	remaining: 45.1ms
25:	total: 46.9ms	remaining: 43.3ms
26:	total: 48.8ms	remaining: 41.5ms
27:	total: 50.5ms	remaining: 39.7ms
28:	total: 52.4ms	remaining: 38ms
29:	total: 54.2ms	remaining: 36.1ms
30:	total: 56ms	remaining: 34.3ms
31:	total: 57.7ms	remaining: 32.5ms
32:	total: 59.4ms	remaining: 30.6ms
33:	total: 61.2ms	remaining: 28.8ms
34:	total: 62.9ms	remaining: 27ms
35:	total: 64.8ms	remaining: 25.2ms
36:	total: 66.6ms	remaining: 23.4ms
37:	total: 68.3ms	remaining: 21.6ms
38:	total: 70.1ms	remaining: 19.8ms
39:	total: 71.8ms	remaining: 18ms
40:	total: 73.7ms	remaining: 16.2ms
41:	total: 75.5ms	remaining: 14.4ms
42:	total: 77.2ms	remaining: 12.6ms
43:	total: 78.9ms	remaining: 10.8ms
44:	total: 80.7ms	remaining: 8.96ms
45:	total: 82.5ms	remaining: 7.17ms

```
46:      total: 84.3ms   remaining: 5.38ms
47:      total: 86.1ms   remaining: 3.59ms
48:      total: 87.8ms   remaining: 1.79ms
49:      total: 89.6ms   remaining: 0us
[Parallel(n_jobs=-1)]: Done 300 out of 300 | elapsed: 5.4min finished
GridSearchCV(cv=StratifiedKFold(n_splits=5, random_state=42, shuffle=True),
             error_score=nan,
             estimator=<catboost.core.CatBoostClassifier object at 0x7f4a152a5950>,
             iid='deprecated', n_jobs=-1,
             param_grid={'iterations': [1, 5, 10, 20, 50],
                          'max_depth': [5, 10, 15, 20],
                          'subsample': [0.66, 0.8, 1]},
             pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
             scoring=None, verbose=1)
```

```
print("Train quality:")
quality(gcv_c.predict(X_train), y_train)
print("\nTest quality:")
quality(gcv_c.predict(X_test), y_test)
```

```
Train quality:
Accuracy: 0.832
Precision: 0.583
Recall: 0.732
F1-score: 0.649
```

```
Test quality:
Accuracy: 0.795
Precision: 0.528
Recall: 0.638
F1-score: 0.578
```

```
plot_roc_curve(gcv_c.predict(X_test), y_test)
```




```
min_child_weight=0.001,
min_split_gain=0.0, n_estimators=100,
n_jobs=-1, num_leaves=31, objective=None,
oob_score=True, random_state=42,
reg_alpha=0.0, reg_lambda=0.0,
silent=True, subsample=1.0,
subsample_for_bin=200000,
subsample_freq=0),
iid='deprecated', n_jobs=-1,
param_grid={'max_depth': [5, 10, 15, 20],
            'max_features': [4, 7, 10, 13],
            'min_samples_leaf': [1, 3, 5, 7]},
pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
scoring=None, verbose=1)
```

```
print("Train quality:")
quality(gcv_l.predict(X_train), y_train)
print("\nTest quality:")
quality(gcv_l.predict(X_test), y_test)
```

```
Train quality:
Accuracy:  0.855
Precision: 0.635
Recall:    0.778
F1-score:  0.699
```

```
Test quality:
Accuracy:  0.791
Precision: 0.512
Recall:    0.631
F1-score:  0.565
```

```
plot_roc_curve(gcv_l.predict(X_test), y_test)
```

▼ Выводы

- без настроек гиперпараметров победила модель XGBOOST auc = 0,707, 2 место - catboost auc = 0,699 (если смотреть по метрике ROC_AUC)
- с настройками гиперпараметров на кросс валидации победила также модель CATBOOST auc = 0,710, 2 место XGBOOST auc = 0,708

Вопрос: не могу понять, почему такое низкое качество выдали мои модели? Дело в настройках фич?

✓ 0s completed at 10:49 PM

